

# Query Processing on Dynamic Networks with Customizable Contraction Hierarchies on Neo4j

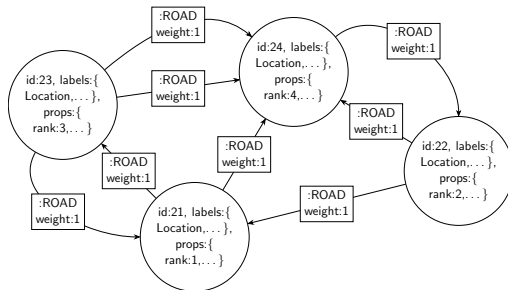
Marius Hahn

Group of Database and Information Systems  
Department of Computer and Information Science  
Faculty of Sciences  
Universität Konstanz

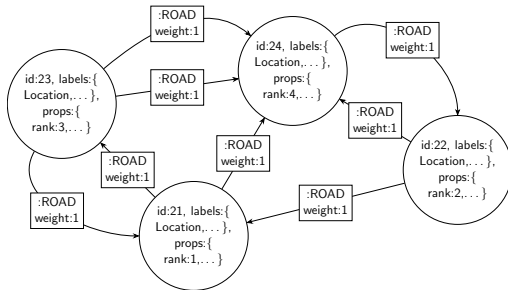
Master Colloquium, 14<sup>th</sup> December 2023

- Graph Databases  $\Rightarrow$  External memory
- Accelerate Shortest Path Queries in Databases
- Why Customizable Contraction Hierarchies?
  - fast for main memory applications
  - reasonable preprocessing time
  - It is updatable
- Test Data  $\Rightarrow$  Road Networks


## Property Graph:




## Property Graph:




## File:



## Property Graph:




## File:




- transformation to data structure with a single dimension

## Property Graph:




## File:




- transformation to data structure with a single dimension
- Databases use HDDs  $\Rightarrow$  slow random access

# Contraction Hierarchies Example




# Contraction Hierarchies Example





# Contraction Hierarchies Example




# Contraction Hierarchies Example

4

# Contraction Hierarchies Example

Let's go from  $v_2$  to  $v_3$




# Contraction Hierarchies Example

Let's go from  $v_2$  to  $v_3$




# Contraction Hierarchies Example

Let's go from  $v_2$  to  $v_3$





# Contraction Hierarchies Example

Let's go from  $v_2$  to  $v_3$




# Customizable Contraction Hierarchies

- CH insert shortcut if shortest path property is violated
- CCH insert shortcut if there is no direct connection




# CCH Weights — Lower Triangles






# CCH Weights — Lower Triangles






# CCH Weights — Lower Triangles



Bottom up for every arc! Also original arcs




# CCH Update — Upper Triangle

- Check if edge(y,z) could rely in edge(x,y)
- If yes  $\Rightarrow$  redo lower triangles for e(y,z)






# CCH Update — Intermediate Triangle

- Check if edge(y,z) could rely in edge(x,y)
- If yes  $\Rightarrow$  redo lower triangles for e(y,z)



# Important vertex not contracted Last


- Contracted Using Edge Difference
- Go from  $v(1)$  to  $v(3)$
- Forward and Backward search are deeper than they should be
- Switch contraction order of  $v(4)$  and  $v(5)$



# Linear Contraction


## 1. linear contraction

- No Shortcuts
- Could happen with ED
- four vertices to expand




## 2. middle vertex first

- Three Shortcuts
- four vertices to expand



## 3. good contraction order

- Three Shortcuts
- four vertices to expand



# Importance Calculation


- add a level  $l(v) = 0$  to each node
- if contracting  $v$ :  $l(v) = \max\{l(v) + 1, l(w)\} \forall w \in N(v)$
- $A(v)$  set of added arcs
- $D(v)$  set of deleted arcs
- $h(a)$  hops an arc represents if unpacked

$$i(v) = l(v) + \frac{|A(v)|}{|D(v)|} + \frac{\sum_{a \in A(v)} h(a)}{\sum_{a \in D(v)} h(a)}$$

## Important theorem

more shortcuts inserted but improves query time!



# Persistence Objectives



- keep only necessary data
  - **rank** → to do the mapping to the input graph
  - arc **weight**
- Store edges that are likely to be request together spacial close
- Use as few space as possible → the less you write the less you read



# Magnetic Disks




- Data is arranged in concentric rings (tracks) on platters
- Tracks are divided into arc-shaped sectors

## One by One

Data is read from and written to disk one **block** at a time

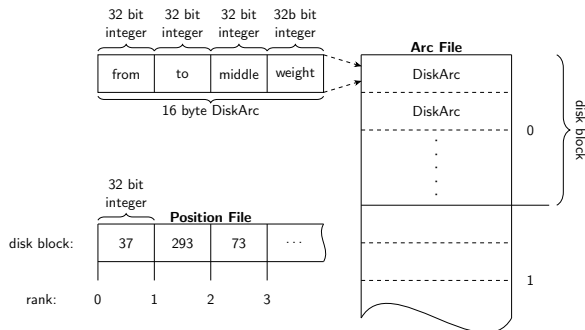
# Transformation to a Table

- Depth-First-Search starting at highest rank
- retrieve only arcs. vertices will be reconstructed from arcs
- remember middle node



start rank	end rank	middle rank	weight
1	3	-1	1
2	3	1	2

# Store Example



- fill all arcs of a vertex into a block
- add block number of rank to position file.  $G_{\uparrow}$  use *from* ;  $G_{\downarrow}$  use *to*
- if next vertex' arcs don't fit anymore  $\rightarrow$  flush block and take next

## Min Block Size

$$\max(d_{\uparrow \max}(v), d_{\downarrow \max}(v)) \leq \frac{\text{diskBlockSize}}{16}$$


# CCH Disk Search (upwards graph example)

1. lazy load vertices  $\Rightarrow$  only start node is loaded without arcs




# CCH Disk Search (upwards graph example)

1. lazy load vertices  $\Rightarrow$  only start node is loaded without arcs
2. settle vertex (right before expanding it)




# CCH Disk Search (upwards graph example)

1. lazy load vertices  $\Rightarrow$  only start node is loaded without arcs
2. settle vertex (right before expanding it)
  - VertexManager requests arc of  $v(x)$  from buffer




# CCH Disk Search (upwards graph example)

1. lazy load vertices  $\Rightarrow$  only start node is loaded without arcs
2. settle vertex (right before expanding it)
  - VertexManager requests arc of  $v(x)$  from buffer
  - Buffer requests arcs from disk if not cached yet



# CCH Disk Search (upwards graph example)


1. lazy load vertices  $\Rightarrow$  only start node is loaded without arcs
2. settle vertex (right before expanding it)
  - VertexManager requests arc of  $v(x)$  from buffer
  - Buffer requests arcs from disk if not cached yet





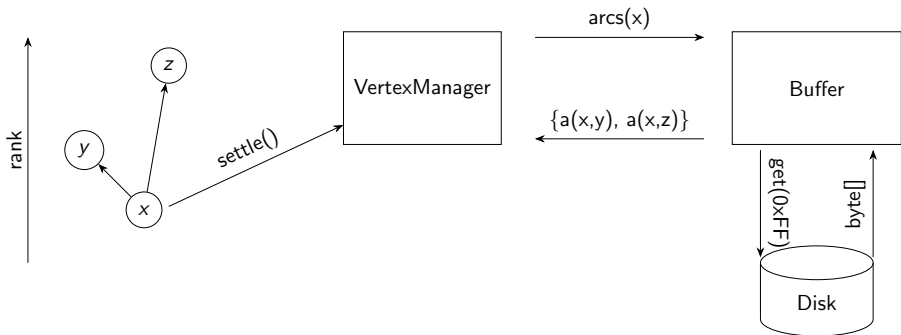
# CCH Disk Search (upwards graph example)

1. lazy load vertices  $\Rightarrow$  only start node is loaded without arcs
2. settle vertex (right before expanding it)
  - VertexManager requests arc of  $v(x)$  from buffer
  - Buffer requests arcs from disk if not cached yet
  - Buffer returns arcs




# CCH Disk Search (upwards graph example)


1. lazy load vertices  $\Rightarrow$  only start node is loaded without arcs
2. settle vertex (right before expanding it)
  - VertexManager requests arc of  $v(x)$  from buffer
  - Buffer requests arcs from disk if not cached yet
  - Buffer returns arcs
  - VertexManager attaches arcs to  $v(x)$




# Circular Buffer



# Circular Buffer



# Circular Buffer



# Circular Buffer

positions:

rank	1	2	$\max(\text{rank})$
position	2	5	-1

DiskArc[]

a(1,x)	a(1,y)	a(1,z)	a(2,x)	a(2,y)	a(2,z)

index:      0            1            2            3            4            5            size

↑  
writePointer

# Circular Buffer

positions:

rank	1	2	3	max(rank)
position	2	5	3	-1

DiskArc[]

a(3,x)	a(3,y)	a(3,z)	a(3,zx)	a(2,y)	a(2,z)

index:      0            1            2            3            4            5            size

remove incomplete edge set from position

writePointer  
↑

# Circular Buffer

positions:

rank	1	3	$\max(\text{rank})$
position	2	3	-1

DiskArc[]

a(3,x)	a(3,y)	a(3,z)	a(3,zx)	a(2,y)	a(2,z)

index:      0            1            2            3            4            5            size

↑  
writePointer



# Circular Buffer

positions:

rank	1	3	max(rank)
position	2	3	-1

DiskArc[]

a(3,x)	a(3,y)	a(3,z)	a(3,zx)	a(2,y)	a(2,z)

index:      0            1            2            3            4            5            size

- Retrieve Arcs  $\Rightarrow$  iterate backwards from position until start vertex differs
- If arc is doesn't start with requested rank  $\Rightarrow$  remove position and refetch

# General Results

	New York	Colorado	Florida	California +Nevada	Great Lakes	Eastern USA
$ V $	264,346	435,666	1,070,376	1,890,815	2,758,119	3,598,623
$ A $	730,100	1,042,400	2,687,902	4,630,444	6,794,808	8,708,058
$ S $	2,153,002	1,680,290	4,397,804	8,598,552	17,833,050	17,712,722
$\frac{ S }{ A }$	2.93	1.59	1.62	1.85	2.62	2.03
$t_{contraction}$	545 s	233 s	579 s	4,384 s	25.29 h	23.29 h
$\max(d(v))$	1,150	629	785	1,252	2,433	2,391
$ \bigcirc A_{\uparrow} $	23.1 MB	21.9 MB	56.9 MB	107 MB	201 MB	215 MB
pos-file $_{\uparrow}$	1.1 MB	1.8 MB	1.3MB	7.6 MB	11.1 MB	14.4 MB
$t_{dijkstra}$	0.816 s	0.549 s	2.630 s	4.858 s	5.425 s	5.387 s
$t_{cch}^{640kB}$	0.140 s	0.122 s	0.147 s	0.289 s	0.732 s	0.727 s
$I/O^{640kB}$	574	437	500	899	1671	1572
$t_{update}^{640kB}$	90 s	51 s	142 s	444 s	1827s	1557s
$t_{cch-updated}^{640kB}$	0.147 s	0.129 s	0.150 s	0.302 s	0.783 s	0.855 s
$I/O_{cch-upd.}^{640kB}$	569	457	516	924	2779	2716
$t_{cch-updated}^{20\%}$	0.136 s	0.130 s	0.092 s	0.183 s	0.660 s	0.680 s
$I/O_{cch-upd.}^{20\%}$	315	307	226	283	804	680
$t_{cch-updated}^{100\%}$	0.062 s	0.038 s	0.039 s	0.099 s	0.438 s	0.479 s

	New York	Colorado	Florida
$ V $	264,346	435,666	1,070,376
$ A $	730,100	1,042,400	2,687,902
$ S $	2,153,002	1,680,290	4,397,804
$\frac{ S }{ A }$	2.93	1.59	1.62
$t_{contraction}$	545 s	233 s	579 s
$\max(d(v))$	1,150	629	785
$ \bigcirc A_{\uparrow} $	23.1 MB	21.9 MB	56.9 MB
pos-file $_{\uparrow}$	1.1 MB	1.8 MB	1.3MB
$t_{update}$	90 s	51 s	142 s

	New York	Colorado	Florida
$t_{dijkstra}$	0.816 s	0.549 s	2.630 s
$t_{cch}^{640kB}$	0.140 s	0.122 s	0.147 s
$t_{cch\text{--}updated}^{640kB}$	0.147 s	0.129 s	0.150 s
$t_{cch\text{--}updated}^{20\%}$	0.136 s	0.130 s	0.092 s
$t_{cch\text{--}updated}^{100\%}$	0.062 s	0.038 s	0.039 s

	New York	Colorado	Florida
$I/O^{640kB}$	574	437	500
$I/O_{cch-upd.}^{640kB}$	569	457	516
$I/O_{cch-upd.}^{20\%}$	315	307	226

	Colorado	Florida	California +Nevada
$ V $	435,666	1,070,376	1,890,815
$ A $	1,042,400	2,687,902	4,630,444
$ S $	1,680,290	4,397,804	8,598,552
$t_{contraction}$	233 s	579 s	4,384 s
$t_{update}$	51 s	142 s	444 s

	Colorado	Florida	California +Nevada
$\frac{ S }{ A }$	1.59	1.62	1.85
$\max(d(v))$	629	785	1,252
$t_{dijkstra}$	0.549 s	2.630 s	4.858 s
$t_{cch}^{640kB}$	0.122 s	0.147 s	0.289 s
$t_{cch-updated}^{640kB}$	0.129 s	0.150 s	0.302 s
$t_{cch-updated}^{20\%}$	0.130 s	0.092 s	0.183 s
$t_{cch-updated}^{100\%}$	0.038 s	0.039 s	0.099 s

	Colorado	Florida	California +Nevada
$I/O^{640kB}$	437	500	899
$I/O_{cch-upd.}^{640kB}$	457	516	924
$I/O_{cch-upd.}^{20\%}$	307	226	283




	California + Nevada	Great Lakes	Eastern USA
$ V $	1,890,815	2,758,119	3,598,623
$ A $	4,630,444	6,794,808	8,708,058
$ S $	8,598,552	17,833,050	17,712,722
$\frac{ S }{ A }$	1.85	2.62	2.03
$t_{contraction}$	4,384 s	25.29 h	23.29 h
$max(d(v))$	1,252	2,433	2,391
$ \bigcirc A_{\uparrow} $	107 MB	201 MB	215 MB
pos-file $_{\uparrow}$	7.6 MB	11.1 MB	14.4 MB
$t_{update}$	444 s	1827s	1557s


	California +Nevada	Great Lakes	Eastern USA
$t_{dijkstra}$	4.858 s	5.425 s	5.387 s
$t_{cch}^{640kB}$	0.289 s	0.732 s	0.727 s
$t_{cch\text{--}updated}^{640kB}$	0.302 s	0.783 s	0.855 s
$t_{cch\text{--}updated}^{20\%}$	0.183 s	0.660 s	0.680 s
$t_{cch\text{--}updated}^{100\%}$	0.099 s	0.438 s	0.479 s

	California +Nevada	Great Lakes	Eastern USA
$I/O^{640kB}$	899	1671	1572
$I/O_{cch-upd.}^{640kB}$	924	2779	2716
$I/O_{cch-upd.}^{20\%}$	283	804	680



# Dijkstra vs. CCH — Query Time




# Dijkstra vs. CCH — Query Time



# Expanded Vertices




# Buffer 640 kB




# Buffer 20% Arc Count

Colorado I/O Circular Buffer



Great Lakes I/O Circular Buffer





# Conclusion

- We can accelerate Graph Databases with CCH
- The major problem is to flatten the graph
- Try it with a Relational Database