

Query Processing on Dynamic Networks with Customizable Contraction Hierarchies on Neo4j


MSc Thesis Title

by

Marius Hahn

at the

Universität
Konstanz



Faculty of Sciences
Department of Computer and Information Science

- | | |
|-----------------|------------------------------------|
| 1. Evaluated by | Prof. Dr. Theodoros Chondrogiannis |
| 2. Evaluated by | Prof. Dr. Sabine Storandt |

Konstanz, 2023

Abstract

Contents

1	Introduction	4
2	Related Work	5
2.1	Algorithmic History	5
2.2	Contraction Hierarchies Database History	5
3	Preliminary	7
3.1	Notation and Expressions	7
4	Customizable Contraction Hierarchies	8
4.1	Contracting and Searching	8
4.2	Difference between CH and CCH	9
4.3	Metric Dependent Vertex Order	9
4.3.1	Important Vertices not contracted last	9
4.3.2	Linear Query Search Space	10
4.4	Vertex importance	11
4.4.1	Suitability of CCH	11
4.4.2	Metric dependent Importance	11
4.5	Update CCH	11
5	Integration in a Neo4j	14
5.1	Index Graph Data Structure	14
5.2	The Mapping	14
5.3	How to Store the Index Graph	15
5.3.1	Persistence Order	16
5.4	Reading Disk Arcs	17
5.4.1	Circular Buffer	17
5.4.2	Least Recently Used Buffer	17
5.5	The Search	17
6	Experiments	19
	Bibliography	25

CHAPTER 1

Introduction

intro

CHAPTER 2

Related Work

As this is mainly a database paper, we want to divide this chapter in two main sections Algorithmic History and Contraction Hierarchies Database History. Algorithmic History that will give some basic overview what has been published regarding index structures to speed up shortest path queries for graphs. Contraction Hierarchies Database History we will try to give an overview of efforts that have been made to make [DSW16, Customizable Contraction Hierarchies] it suitable for graph databases.

2.1 Algorithmic History

[GSSV12, Contraction Hierarchies] or CH is heavily influenced by the idea of the [BFSS07, Transit-Node] approach and as transit node approach itself, CCH is a technique to speed up [Dij59, Dijkstras Algorithm], which is the most basic and robust algorithm to find shortest path in graphs. CH goes back to the diploma thesis of [GSSD08, Geisberger] in 2008. The [BFSS07, Transit-Node] approach tries to find vertices inside the graph that are more important than others. Important in this case means, these are vertices that reside on many shortest paths. This speeds up especially long distance queries, as one only needs to calculate the distance to the the next transit node of the source and target vertex as the shortest paths between the transit or access nodes will be known.

CH goes even further on the idea of having important vertices. It applies an importance to each vertex in the graph a so called rank. Furthermore it adds edges to the graph, so called shortcuts, that preserve the shortest path property of the graph in case a vertex that is contracted resides on a shortest path between others. When querying a shortest path CH uses a modified bidirectional-dijkstra that is restricted to only visit nodes that are of higher importance, or rank, than the its about to expand next. This method is able to retrieve shortest paths of vertices that have a high spacial distance, however, it is rather static. In case a new edge is added or an edge weight is updated, it might be necessary to recontract the whole graph to preserve the shortest path property.

In 2016 [DSW16, Customization Contraction Hierarchies] or CCH was published. The approach is the same, but in CCH shortcuts are not only added if the contraction violates the shortest path property, they are added if there had been a connection between its neighbors through the just contracted vertex and these neighbors do not own a direct connection through an already existing edge. The shortcut weights are later on calculated through the lowers triangle. Additionally the [DSW16, Customization Contraction Hierarchies] provides an update approach that only updates, edges that are affected by a weight change.

2.2 Contraction Hierarchies Database History

There is one bachelor thesis by Nicolai D’Effremo [D’E19, Some text] that has implemented a version on [GSSV12, Contraction Hierarchies] for Neo4j, one of the most used graph databases

of today in 2023. This implementation shows that even in for databases CH is an index structure worth pursuing, as there was a tremendous speedup of shortest path queries paired with a reasonable preprocessing time. [Zic21] showed in his bachelor thesis that it is even possible to restrict these queries with label constraints. Although CH and CCH have little difference, sadly we could not use much of the code provided by there works. It was deeply integrated into the Neo4j-Platform and since then two major release updates happened that have breaking changes which make it nearly impossible to reuse any of this code.

Finally there is [SSV, Mobile Route Planning] by Peter Sanders, Dominik Schultes, and Christian Vetter. In this paper it is described how one can efficiently store the a CH index structure on a hard drive. It states an interesting technique to how store edge that are likely to be read sequentially spatially close on the hard drive which makes read operations that have to be done during query time fast. The motivation of [SSV, Mobile Route Planning] through was slightly different. They came up with this idea because computation power on mobile devices is limited, so they could precalculate the CH index on a server and then later distribute it to a mobile device.

We will use parts of this idea and partly port it to our database context as we suppose there are many similarities.

CHAPTER 3

Preliminary

As the target platform for this work is the graph database neo4J, we will mostly consider *directed* graphs. From the terminology we always refer *arcs*, which is an directed edge. In some cases we will refer to *edges*, in these cases the direction doesn't play a role.

3.1 Notation and Expressions

We denote a graph $G(V, A)$ in case we mean an *directed* graph, where v is a vertex contained in the vertices $v \in V$ and a is an arc $a \in A$. An arc is uniquely defined by two vertices v_a and v_b such that $v_a \neq v_b$, so there are no loops nor multi edges. An edge additionally has a weight function $w : A \rightarrow \mathbb{R}_{>0}$ it's weight which must be a positive.

We use A as the arc set and a a single arc which is directed. $a \in A$ can be replaced with $e \in E$ which refers to edges that are *undirected*.

G represents the input graph. The contraction graph $G'(V', A')$ is the graph that will be used at contraction for initially building the CCH index structure. A vertex v in will never be really deleted. Instead the rank property $r(v)$ is set to mark this as an already contracted. So $V \equiv V'$ but $A \subseteq A'$ there will be edges added while building the CCH index. $S = A' \setminus A$ is the shortcut set that is added throughout the contraction.

$G^*(V^*, A^*)$ is the search graph while doing one a shortest path query. Furthermore one query will have two search graphs. G_\uparrow^* representing the upwards search graph and the G_\downarrow^* .

Finally there will be the edge set of edges that are written to the disk. These will $\bigcirc A$ will be separated into two sets $\bigcirc A_\downarrow$ and $\bigcirc A_\uparrow$, too.

Customizable Contraction Hierarchies

In this section we will present the basic idea of [DSW16, Customization Contraction Hierarchies] and also work out the main difference between CCH and [GSSV12, Contraction Hierarchies]. It is far from being complete, but there will be some easy examples to show the concept.

4.1 Contracting and Searching

In Figure 1 you can see a contracted graph $G'(V, E')$ on the left. The solid lines represent the original edges E of a graph G . The dashed lines between vertices are shortcuts S that have been added while creating the CCH index graph $G'(V, E')$. The numbers inside the vertices reflect the contraction order.

Contracting a vertex means deleting it. While contracting a vertex we want to preserve its via connection. If a vertex that is contracted resides on a simple path between two vertices of higher rank, and there is no edge $e \in E'$ between these vertices a shortcut has to be inserted between the two. Let's reconstruct the contraction of Figure 1. At first vertex $v(1)$ is removed. As $v(1)$ resides on a simple path between $v(3)$ and $v(5)$ and there is no edge $e(v(3), v(5)) \notin E'$, there must be a shortcut added to keep the via path. The same applies after contracting $v(2)$ for the vertices $v(4)$ and $v(5)$. For all the other vertices we do not need to insert shortcuts.

As we preserved all via paths during the contraction the shortest path can be retrieved by a bidirectional Dijkstra that is restricted such that it only expands vertices of higher rank. Therefore if one wants to retrieve the shortest path between $v(3)$ and $v(4)$ there will be a forward search from $v(3)$ and a backward search from $v(4)$. As we restrict these searches to expand only vertices of higher rank, the only vertices to expand are the start and target vertex. Both will find only one vertex $v(5)$, the highest vertex and the meeting point, too. Finding at least one meeting point in the forward and backward search means there exist a path between them. After merging these paths at the middle vertex $v(5)$ one will obtain the shortest path.

For an arbitrary contracted graph is it possible that there are more than one meeting point. As merging two shortest paths will not necessary lead to an other shortest path, one has to merge all possible meeting points and take the path among the merged ones which has the smallest distance.

The stopping condition for such a CH-Search is either, both forward and backward search, have reached the top vertex so there is no further vertex to expand, which happens in the example of figure 1 or, backward and forward search exceed the length that has already been found among the merged paths.




Figure 1 The numbers inside the vertices represent their contraction order




Figure 2 The left represents a CH and the right a CCH contracted graph

4.2 Difference between CH and CCH

Looking at the left graph in Figure 2 it has been contracted in the CH way, whereas the right is the CCH way. We explicitly state this here because we have found paper [OYQ⁺20] that mix up these well known names, claiming they to Contraction Hierarchies CH while actually doing Customizable Contraction Hierarchies CCH. The main difference is, CH will only insert an shortcut between two vertices if the vertex that is contracted resides on the shortest path between two of its neighbors. When vertex $v(1)$ is contracted there is no shortcut inserted as vertex $v(1)$ is not on the shortest path between which is via vertex $v(4)$.

Whereas in the CCH case the edge weights do not play a role a contraction time. If a vertex is contracted and there is no direct connection between two of its neighbors, one has to insert a shortcut. This gives the advantage that later on we can easily update edge weights without inserting new shortcut, as all possibly needed shortcuts already exist.

Let's complete this example by updating the edge $e(v(2), v(4))$ that currently has the weight of $w(e) = 1$ to $w(e) = 5$. Now the vertex $v(1)$ is on the shortest path between vertex $v(2)$ and $v(3)$. To update the CH graph we have to insert an edge between vertex $v(2)$ and $v(3)$ whereas the topological structure of the CCH remains the same, one only need to update the weight and the middle vertex of the already give shortcut edge.

4.3 Metric Dependent Vertex Order

There are two ways to get a suitable vertex order. A so called *metric independent* and a so called *metric dependent* one. The metric independent recursively uses balanced separator to determine a vertex ordering[DSW16]. Although this is the superior method, it is not used in this paper writing an algorithm that calculates balanced separators isn't trivial, and we are not aiming for optimizing the contraction process. The metric dependent order mainly uses the edge difference ED to determine which vertex is to be contracted next. The ED is determined as the $|edgesToInsert| - |edgesToRemove|$. The fewer edges are inserted during contraction the fewer edges will be contained by the final graph, therefore fewer edges to expand in a search. However using only the edge differences doesn't lead to desired result. This is because during contraction there will be areas that get less dense than others. There are two problems that can arise. One is that important vertices are not contracted last. The other is the search space of the query gets linear although it could be logarithmic.

4.3.1 Important Vertices not contracted last

Looking at figure 3, this is a possible contraction order, if only the ED is used to contract vertices. At the beginning the vertices with rank 1, 2, 3, 5 have the same edge difference, which is




Figure 3 The numbers inside the vertices represent their contraction order




Figure 4 Linear Contraction

$ED = -1$. Vertex after vertex is removed and no shortcut is inserted. This happens until there are only $v(4)$ and $v(5)$ left. Now $v(4)$ has an $ED = -1$, too, same as vertex 5. Therefore the algorithm contracts $v(4)$ before $v(5)$.

However this is not the desired result. There are six $e(v(1), v(2))$, $e(v(1), v(3))$, $e(v(1), v(5))$, $e(v(2), v(3))$, $e(v(2), v(5))$, $e(v(3), v(5))$ shortest paths that involve $v(4)$, all the other vertices do not encode any shortest path, so $v(4)$ should be contracted last. The search graph on the right of Figure 3 shows why. Imagine we do a shortest path query between $v(1)$ and $v(3)$. After expanding both, the forward and the backward search to $v(4)$, there is yet another vertex we'll have to expand $v(5)$. Although as you can see in the original graph on the right, its not possible that $v(5)$ is on the shortest path. Therefore a better contraction order would be a in Figure 1. This can be overcome by the method that is explained in section 4.4.

4.3.2 Linear Query Search Space

Regarding figure 4 there are three possible index graphs G' of one and the same base graph G . The numbers inside the vertices represent the contraction order.

The first one could be contracted using the edge difference ED , as always one of the outer vertices with $ED = -1$ was contracted. On the one hand it reaches the optimum in case for *least shortcuts inserted*. On the other though it has the worst search space among the three vertex orderings. To get from vertex $v(1)$ to $v(5)$ we have to expand four vertices.

The second G' one contracts the middle vertices, which encodes the most shortest paths, first and therefore inserts three shortcuts. Although this example has a lot of shortcuts, there are still a lot of vertices to expand in some cases. In case every vertex of G has a weight of 1, and one wants to go from $v(1)$ to $v(5)$ the forward search will have to expand four vertices as in the upper first example.

The third example contracts the middle vertex last. At first it contracts the vertices right next to the middle vertex. Therefore we have to insert shortcuts between $e(v(3), v(5))$ and $e(v(4), v(5))$, so no matter what source, target pair we are trying to find in this example, the forward and the backward search will have to expand at most one single vertex. This example additionally

shows that recursively finding a balanced separator, as proposed in [DSW16, Customization Contraction Hierarchies], is very a promising method to obtain a good contraction order.

4.4 Vertex importance

As shown in section 4.3.1 and 4.3.2 there are vertices that are more important than other vertices. Contracting these vertices late is key to get a efficient search later on.

4.4.1 Suitability of CCH

As it is important to contract important vertices last, the advantage one gets making a CCH search over a simple dijkstra run depends whether the base graph $G(V, E)$ has vertices that are more important than others. A vertex $v \in E$ is important if there are many shortest paths that contain this very vertex. Therefore if it is possible to calculate a small balanced separator on G , CCH will be able to show its whole advantage. To dive deeper into this topic, have a look at [BS20, Lower Bounds and Approximation Algorithms for Search Space Sizes in Contraction Hierarchies].

4.4.2 Metric dependent Importance

As shown above, taking only the edge difference ED into account doesn't necessarily lead to a proper order, we decided to take the vertex importance calculation that is proposed by [DSW16, Customization Contraction Hierarchies]. To every vertex we add the level property $l(v)$. The level of the vertex is initially set to 0. If a neighbor $w = N(v)$ is contracted level is set to $l(v) = \max\{l(v) + 1, l(w)\}$. For every arc $a \in A'$ we add the hop length to the arc $h(a)$. The hop length is equal to the number of arcs, this arc represents when fully unpacked. Additionally, we denote as $A(v)$ the set of inserted arcs after the contraction of v and $D(v)$ the set of removed arcs. We calculate the importance $i(v)$ as follows:

$$i(v) = l(v) + \frac{|A(v)|}{|D(v)|} + \frac{\sum_{a \in A(v)} h(a)}{\sum_{a \in D(v)} h(a)}$$

Our tests show that this importance calculation results in a slightly increase in the amount of shortcuts added, but the maximum Vertex degree is smaller. Which speeds up the contraction process towards the end. Additionally the average search time decreases as the search space decreases too.

4.5 Update CCH

The biggest advantage of CCH over CH is, that it is easy to update without the need of changing the topological structure of the index graph. This is the reason why CCH can be interesting for graph databases. If an arc's weight $w(a(x, y))$ increases or decreases this can result in a weight change on arcs that connect vertices of higher rank than x, y . We determine all arcs of the input graph G that have been changed and push them to a priority queue. The queue always pops the arc $a(x, y)$ with the lowest rank of the start vertex x . If there are multiple it pops the one with the lowest rank of y among the ones with the lowest rank of x . Then we determine the new weight of the arc using the lower triangles. If there is a lower triangle that can be used as a pass

Algorithm 1 Update

```
1: procedure update()(G)
2:    $Q \leftarrow G'.updatedEdges(G)$ ;
3:   while  $Q \neq \emptyset$  do
4:      $a \leftarrow Q.poll()$ ;
5:      $oldWeight \leftarrow w(a)$ 
6:      $newWeight \leftarrow determineNewWeight(a)$ 
7:     if  $oldWeight \neq newWeight$  then
8:        $w(a) \leftarrow newWeight$ 
9:        $checkTriangles(Q, oldWeight, upperTriangles(a))$ 
10:       $checkTriangles(Q, oldWeight, intermediateTriangles(a))$ 
11:     end if
12:   end while
13: end procedure
14: procedure checkTriangles( $Q, oldWeight, triangles$ )
15:   for all  $triangle$  in  $triangles$  do
16:     if  $triangles.c() == triangle.b() + oldWeight$  then
17:        $Q.push(triangle.c())$ 
18:     end if
19:   end for
20:   return  $triangles$ 
21: end procedure
```

through such that the arc weight in G' does not change we do nothing. If the weight of the arc has changed we assign the new weight to the arc. Then we check all upper triangles, as drawn in figure 5, of $a(x, y)$ if there is an upper arc; denoted by c in figure 5, that is influenced by this very change. If it is influenced by this change we push it to the priority queue. We do the same with all intermediate triangles.




Figure 5 Update Triangles

Algorithm 2 Compute Triangles

```
1: procedure updateTriangles(arc, neighbors, predicate)
2:   triangles  $\leftarrow \{\}$ ;  $x \leftarrow \min(\text{arc.start}, \text{arc.end})$ ;  $y \leftarrow \max(\text{arc.start}, \text{arc.end})$ ;
3:   for all  $z$  in neighbors do
4:     if predicate then
5:       if upwards(arc) then triangles.add(Triangle( $a(x, y)$ ,  $a(z, x)$ ,  $a(z, y)$ ))
6:       else triangles.add(Triangle( $a(x, y)$ ,  $a(x, z)$ ,  $a(y, z)$ ))
7:       end if
8:     end if
9:   end for
10:  return triangles
11: end procedure
12: procedure intermediateTriangles(arc)
13:  neighbors  $\leftarrow N_{\uparrow}(\text{arc.start}) \cap N_{\downarrow}(\text{arc.end})$ 
14:  return updateTriangles(arc, neighbors,  $x < z < y$ )
15: end procedure
16: procedure upperTriangles(arc, neighbors)
17:  neighbors  $\leftarrow N_{\uparrow}(\text{arc.start}) \cap N_{\downarrow}(\text{arc.end})$ 
18:  return updateTriangles(arc, neighbors,  $x < y < z$ )
19: end procedure
```

CHAPTER 5

Integration in a Neo4j

In this section it is described how "Customizable Contraction Hierarchies" CCH is integrated into Neo4j. CCH arguments the input graph, which means it inserts arcs, so called shortcuts, that do not belong to the original data. To keep the change to the input graphs as little as possible we decided to not insert any arc into the graph that is stored inside the neo4j database, but introduce another graph data structure, the index graph. This index graph has an mapping to the input graph that is held by the database, by inserting two properties into the node of the input graph. The *rank* this vertex has in the index graph and the *indexing weight* it had during the last customization process. This gives yet another two advantages. One is that we get full control about the graph representation which is helpful to efficiently store and read the index graph for the disk. Another is that the with this approach it makes it easier to later on port the idea to another graph database manufactures.

5.1 Index Graph Data Structure

The index graph data structure is neither a adjacency list nor adjacency matrix. There is a vertex object that has two hash tables. One for incoming arc and one for outgoing arcs. The hash tables keys are of type vertex and the value is the arc. An arc has a reference to its start vertex and one to its end vertex.

A disadvantage of this model could be that some modern hardware optimization that exist for arrays do not match with this data structure. When using an array, the values this array are stored sequentially in main memory. When one value of an array is accessed by the CPU, modern hardware reads subsequent values into the CPU-cache because it is likely that they are accessed right after it. The model of the index graph is a linked data structure, a bit like a linked list. The elements of an linked list are contained somewhere in main memory. There is no guarantee that subsequent values have any spacial proximity. Therefore the just explained hardware optimization will not give any advantage.

However, this makes the makes the graph traversal easy. Additional it makes it very efficient to explore the neighborhood of a vertex. There is no array traversal to find a vertex and only one hash table lookup for finding an arc of a vertex. Additionally these hash tables only contain few elements. This makes this data structure efficient anyway. Test on small graphs [Oldenburg] show that cch queries can be answered in less than one millisecond, which is close to what we tested with the original cch application.

5.2 The Mapping

The in memory data structure of neo4j is similar to the just explained index graph data structure in section 5.1. A *node* has a collection of *relationships* and a *relationship* has a reference to its *start node* and *end node*. As neo4j is a full blown property graph nodes and relationship contain a lot of other information. A node has a collection of *labels*, relationship has a *type*. The class




Figure 6 mapping

Node and the class *Relationship* are both derived from the class *Entity* which also has a collection of properties as well as an *id* that is managed by the database system. Note that, as of version Neo4j 5.X, this *id* can change over time and should not be used to make mappings to external systems. Additionally worth to mentioning here is that the Neo4j system shifted its *id* concept as it moved from major release 4 to 5. Until major release 4 every entity had a unique integer identifier. Since major release 5 every entity has a string identifier which is a UUID and the old *id* identifier isn't guaranteed to be unique anymore. It is deprecated and marked for removal.

As just explained there are a lot of information in this data structure. A lot of information we don't need. Looking at 6 we only want to keep track of the information that is needed for the CCH index. Additionally as disks are divided into blocks and sectors we want to flatten the graph which is in memory more looks like a tree to a structure that looks like a table. Therefore we decided that the disk data structure only consists of edges $\bigcirc A$. A disk edge $a \in \bigcirc A$ consists of four values, the *start rank*, the *end rank*, the *start rank* and the *weight*. The middle node is set -1 in case that this arc, is an arc of the input graph. We will get two edge sets $\bigcirc A_{\downarrow}$ for the downwards graph and $\bigcirc A_{\uparrow}$ upwards graph. $\bigcirc A_{\downarrow}$ contains all downward edge that which are needed for the backward search and $\bigcirc A_{\uparrow}$ contains all upwards arcs that are needed for the forward search.

During the contraction every node gets a rank assigned. This rank is the only change that is made to the Neo4j data structure and its the mapping identifier between the input graph G and the index graph G' . G' will then be used to generate $\bigcirc A_{\downarrow}$ and $\bigcirc A_{\uparrow}$.

5.3 How to Store the Index Graph

After generating the disk arc sets $\bigcirc A_{\downarrow}$ and $\bigcirc A_{\uparrow}$, we now want to store them as efficiently as possible to the disk. To refine the definition of a disk arc. It consist of four values *start rank*, the *end rank*, the *start rank* and the *weight*. The first three are 32 bit signed integer, which gives a maximum indexable amount of vertices for 2^{16} . The last one, the arc *weight* is a 32 bit

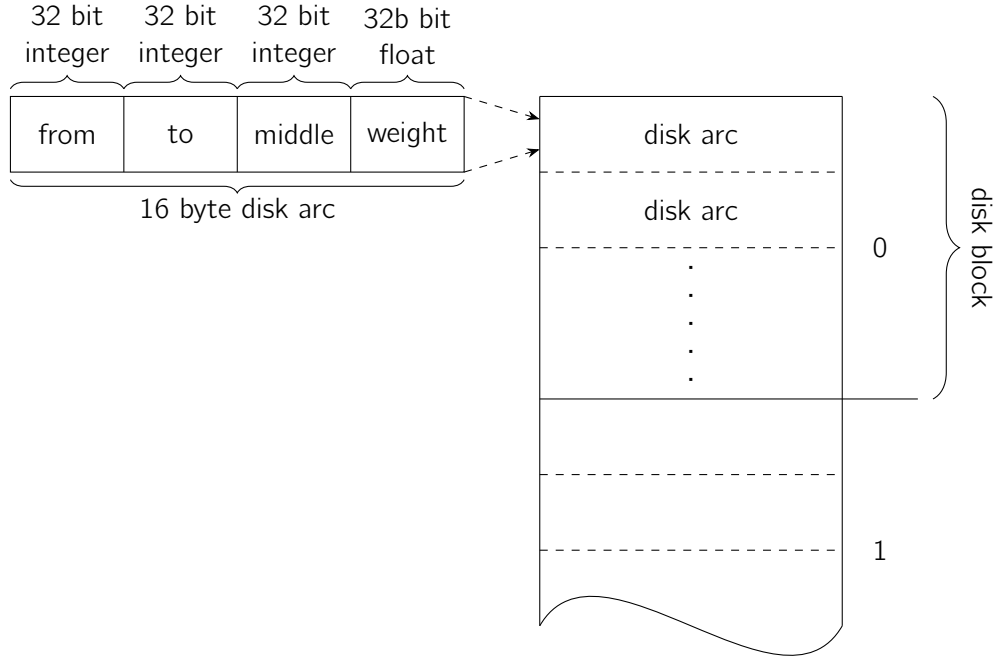


Figure 7 Disk Block

floating point number. One could argue that 32 bit are not very precise, though our experiments we never had any imprecision problem. Furthermore, little imprecision would not cause a big problem, as the index graph is only needed to find the shortest path. The exact weight can later on be retrieved after the shortest path is resolved in G .

This 16 Byte disk collected to disk blocks, that have at least the size of one disk block on the systems hard drive, as usually the system will always read a complete disk block even if you request only 16 Byte of information. But it is possible to make blocks bigger than one disk block. This can be wanted if you have a big read cache, or even need it the highest outgoing or in ingoing arc degree exceed the $\frac{diskBlockSize}{16}$ as you will see in 5.3.1.

5.3.1 Persistence Order

As the disk arc are later on is used for the CCH search, we want to sequentially write them in a way that provides a high spatial proximity of vertices that are likely to get requested together. Here we will adopt the idea of [SSV, Mobile Route Planning]. In the transformation from G' to its disk arcs $\bigcirc A_{\uparrow}$ and $\bigcirc A_{\downarrow}$ we do a simple depth first search on all ingoing arcs on the target rank to determine the order for $\bigcirc A_{\uparrow}$. We do the same for all outgoing arcs on the source rank to determine the order for $\bigcirc A_{\downarrow}$. There is one file for storing all upward arc and on file storing all downward arcs. Each of this files has a position file that belongs to it, as shown in figure 8. After every vertex that is expanded, all its arcs are pushed to a buffer of size disk block. If the buffer is full or the number of arcs doesn't fit anymore, the buffer is flushed to the current position in the arc file and the file position is incremented by 1. The store position is saved in the position file of the arc file, such that we can find the arcs back later on. If the buffer wasn't complete full all remaining arc slots are filled with dummy arcs that contain only -1 at every property.




Figure 8 Position File

5.4 Reading Disk Arcs

If one wants to get all upwards arcs of rank i one need take the upwards position file, retrieve the integer j that is stored at index i , and then read the complete block j in the upwards arc file. There There one will get an array, contain the requested arcs but also some other. These arcs are likely to be request next. Therefore we want to keep them in memory. We implemented two buffer a *circular buffer* and a *least recently used buffer* LRU

5.4.1 Circular Buffer

For the circular buffer we simply used an array of disk arcs. If we reach the end of the buffer we restart overwriting the values from the array start. To get the all arcs of a rank one request that rank number. There is a position hash table which tells the start position of that rank inside the buffer. If it is missing, the containing disk block is read to the buffer. It continues to read sequentially until the request rank and the read arc doesn't belong together anymore. This buffer has the advantage that we can exactly determine the amount of arcs we buffering. Also as it is just a simple array, it will be easy for the operation system to cache it. The disadvantage is that it is possible that we request arc sets very often as it is possible they get evicted just before request again.

5.4.2 Least Recently Used Buffer

Cache is a *java.util.LinkedHashMap*. This class provides the possibility to evicted the entry that has been requested longest time ago. In our case it maps ranks to sets of disk arcs. We can only determine how many disk arc sets we have in memory and disk arc sets do not have always the same size. Higher rank vertex usually have bigger sets as they are of higher degree. The advantage is, it is very easy to implement and therefore very resilient to programming errors.


5.5 The Search


The search brings all things explained in this chapter together.

At the beginning there are to index graphs initialized the upwards graph $G'_\uparrow(V_\uparrow, A_\uparrow)$ and the downwards graph $G'_\downarrow(V_\downarrow, A_\downarrow)$. We are looking for the shortest path from the source vertex $v(s)$ to the target vertex $v(t)$. The vertex set V_\uparrow of the upwards graph $G'_\uparrow(V, A_\uparrow)$ only contains one vertex $v(s)$ and the A_\uparrow only contains the upward edges to $v(s)$. The vertex set V_\downarrow of the downwards graph $G'_\downarrow(V_\downarrow, A_\downarrow)$ only contains the $v(t)$ and the edge set A_\downarrow only contain the arcs to $v(t)$. Now both G'_\uparrow and G'_\downarrow are alternately expanded with CH-Dijkstra. When the CH-Dijkstra has


decided which vertex to expand next, the buffer will be requested to load its neighborhood. This can be compared to Command & Conquer or a lot of other strategy real-time strategy video games where you start at a map that is almost completely grey and the map is only load to the position you are plus some padding. The rest is just a plain CH-Search where we stop after we determined the correct shortest path.


Experiments






Disk access over Expand Nodes in CCH





Disk access over Expand Nodes in CCH



	New York	Colorado	Florida	California + Nevada
Verticies	264,346	435,666	1,070,376	1,890,815
Edges	733,846	1,057,066	2,712,798	4,657,742
Shortcuts	2,153,002	1,680,290	4,397,804	8,598,552
Shortcuts added Relative	2.93	1.59	1.62	1.85
Contraction Time	545 s	233 s	579 s	4,384 s
Max In Degree	1,150	629	785	1,252
Out Arcs Disk Size	23.1 MB	21.9 MB	56.9 MB	106.5 MB
Out Position File Size	1.1 MB	1.8MB	1.3MB	7.6 MB
In Arcs Disk Size	23.1 MB	21.9 MB	56.9 MB	106.5 MB
In Position File Size	1.1 MB	1.8MB	1.3MB	7.6 MB
Block filling	0.99	0.95	-122.0	-122.0
AVG Dijkstra Search Time	-122.0	-122.0	-122.0	-122.0
AVG Search Time 40960	-122.0	-122.0	-122.0	-122.0
AVG Search Time 25%	-122.0	-122.0	-122.0	-122.0
AVG Search Time 50%	-122.0	-122.0	-122.0	-122.0
AVG Search Time 75%	-122.0	-122.0	-122.0	-122.0
AVG Search Time 100%	-122.0	-122.0	-122.0	-122.0
AVG Disk Access 40960	-122.0	-122.0	-122.0	-122.0
AVG Disk Access 25%	-122.0	-122.0	-122.0	-122.0
AVG Disk Access 50%	-122.0	-122.0	-122.0	-122.0
AVG Disk Access 75%	-122.0	-122.0	-122.0	-122.0
Update Time	-122.0	-122.0	-122.0	-122.0
AVG Search Time 40960 after multi update	-122.0	-122.0	-122.0	-122.0
AVG Disk Access 40960 after multi update	-122.0	-122.0	-122.0	-122.0

Table 1 thisTable.

Bibliography

- [BFSS07] Holger Bast, Stefan Funke, Peter Sanders, and Dominik Schultes. Fast routing in road networks with transit nodes. *Science*, 316(5824):566–566, apr 2007.
- [BS20] Johannes Blum and Sabine Storandt. Lower bounds and approximation algorithms for search space sizes in contraction hierarchies. 2020.
- [D’E19] Nicolai D’Effremo. An external memory implementation of contraction hierarchies using independent sets, 2019.
- [Dij59] E. W. Dijkstra. A note on two problems in connexion with graphs. *Numerische Mathematik*, 1(1):269–271, dec 1959.
- [DSW16] Julian Dibbelt, Ben Strasser, and Dorothea Wagner. Customizable contraction hierarchies. *ACM Journal of Experimental Algorithmics*, 21:1–49, apr 2016.
- [GSSD08] Robert Geisberger, Peter Sanders, Dominik Schultes, and Daniel Delling. Contraction hierarchies: Faster and simpler hierarchical routing in road networks. pages 319–333, 2008.
- [GSSV12] Robert Geisberger, Peter Sanders, Dominik Schultes, and Christian Vetter. Exact routing in large road networks using contraction hierarchies. *Transportation Science*, 46(3):388–404, aug 2012.
- [OYQ⁺20] Dian Ouyang, Long Yuan, Lu Qin, Lijun Chang, Ying Zhang, and Xuemin Lin. Efficient shortest path index maintenance on dynamic road networks with theoretical guarantees. *Proceedings of the VLDB Endowment*, 13(5):602–615, jan 2020.
- [SSV] Peter Sanders, Dominik Schultes, and Christian Vetter. Mobile route planning. pages 732–743.
- [Zic21] Anton Zickenberg. A contraction hierarchies-based index for regular path queries on graph databases, 2021.